

KerasNeuralNetwork

September 29, 2021

1 Import packages

```
[2]: # import packages
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
# packages for building the model
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Activation, Dense
from tensorflow.keras import losses
from tensorflow.keras import metrics
# packages for training model
from tensorflow.keras import optimizers
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.metrics import categorical_crossentropy
# misc package
import category_encoders as category_encoder
```

2 Import and view raw data

```
[3]: data = pd.read_csv('heart.csv')
```

```
[4]: data.columns
```

```
[4]: Index(['Age', 'Sex', 'ChestPainType', 'RestingBP', 'Cholesterol', 'FastingBS',
          'RestingECG', 'MaxHR', 'ExerciseAngina', 'Oldpeak', 'ST_Slope',
          'HeartDisease'],
          dtype='object')
```

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	918 non-null	int64
1	Sex	918 non-null	object
2	ChestPainType	918 non-null	object
3	RestingBP	918 non-null	int64
4	Cholesterol	918 non-null	int64
5	FastingBS	918 non-null	int64
6	RestingECG	918 non-null	object
7	MaxHR	918 non-null	int64
8	ExerciseAngina	918 non-null	object
9	Oldpeak	918 non-null	float64
10	ST_Slope	918 non-null	object
11	HeartDisease	918 non-null	int64

dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB

```
[6]: # 918 samples, 12 features
data.shape
```

```
[6]: (918, 12)
```

3 Preprocessing of Data

- Check for null values
- Check for duplicated values
- convert categorical data to numerical data
- Normalize data
- Split data into train and test samples/labels
- Convert datasets to numpy arrays to be able to pass data to Keras models

3.1 Preprocessing data: Check for null and/or duplicated values

```
[7]: # check for missing values - there is no missing data
data.isnull().sum(axis=0)
```

```
[7]: Age          0
     Sex          0
     ChestPainType  0
     RestingBP     0
     Cholesterol   0
     FastingBS     0
     RestingECG    0
     MaxHR         0
     ExerciseAngina 0
     Oldpeak       0
     ST_Slope      0
```

```
HeartDisease      0
dtype: int64
```

```
[8]: # check for duplicate values - there is no duplicated data
data.duplicated().sum(axis=0)
```

```
[8]: 0
```

3.2 Preprocessing data: View unique values for categorical features

```
[9]: category_columns = []

# viewing the unique values, number of dimensions, and shape of each column in
↳ the data frame
for col in data.columns:
    if data[col].dtype == 'object':
        print(f'{col}')
        print(f'Values: {data[col].unique()}')
        print('\n')
        category_columns.append(col)
print(f'The following categories: {category_columns} shall be converted to
↳ numerical data via pd.get_dummies')
```

```
Sex
Values: ['M' 'F']
```

```
ChestPainType
Values: ['ATA' 'NAP' 'ASY' 'TA']
```

```
RestingECG
Values: ['Normal' 'ST' 'LVH']
```

```
ExerciseAngina
Values: ['N' 'Y']
```

```
ST_Slope
Values: ['Up' 'Flat' 'Down']
```

```
The following categories: ['Sex', 'ChestPainType', 'RestingECG',
'ExerciseAngina', 'ST_Slope'] shall be converted to numerical data via
pd.get_dummies
```

```
[10]: # labels - 0 does not have heart disease, 1 does have heart disease
data['HeartDisease'].unique()
```

```
[10]: array([0, 1])
```

3.3 Preprocessing data: Converting categorical values into numerical values

```
[11]: # Convert the nominal categorical label values to numerical values

# get_dummies method (one hot encoding technique for multi-value categorical
    ↪ variables)
mod_data = data.copy()
mod_data = pd.get_dummies(mod_data, columns=category_columns)
mod_data.head(2)
```

```
[11]:   Age  RestingBP  Cholesterol  FastingBS  MaxHR  Oldpeak  HeartDisease  \
0   40         140          289          0    172        0.0           0
1   49         160          180          0    156        1.0           1

   Sex_F  Sex_M  ChestPainType_ASY  ...  ChestPainType_NAP  ChestPainType_TA  \
0       0      1                0  ...                0                0
1       1      0                0  ...                1                0

   RestingECG_LVH  RestingECG_Normal  RestingECG_ST  ExerciseAngina_N  \
0                0                  1                0                1
1                0                  1                0                1

   ExerciseAngina_Y  ST_Slope_Down  ST_Slope_Flat  ST_Slope_Up
0                0                0                0                1
1                0                0                1                0

[2 rows x 21 columns]
```

```
[12]: mod_data.columns
```

```
[12]: Index(['Age', 'RestingBP', 'Cholesterol', 'FastingBS', 'MaxHR', 'Oldpeak',
        'HeartDisease', 'Sex_F', 'Sex_M', 'ChestPainType_ASY',
        'ChestPainType_ATA', 'ChestPainType_NAP', 'ChestPainType_TA',
        'RestingECG_LVH', 'RestingECG_Normal', 'RestingECG_ST',
        'ExerciseAngina_N', 'ExerciseAngina_Y', 'ST_Slope_Down',
        'ST_Slope_Flat', 'ST_Slope_Up'],
        dtype='object')
```

```
[13]: # separating features from output (HeartDisease)
features = [col for col in mod_data.columns if col != 'HeartDisease']
output = ['HeartDisease']
features, len(features)
```

```
[13]: ([ 'Age',
        'RestingBP',
        'Cholesterol',
        'FastingBS',
        'MaxHR',
        'Oldpeak',
        'Sex_F',
        'Sex_M',
        'ChestPainType_ASY',
        'ChestPainType_ATA',
        'ChestPainType_NAP',
        'ChestPainType_TA',
        'RestingECG_LVH',
        'RestingECG_Normal',
        'RestingECG_ST',
        'ExerciseAngina_N',
        'ExerciseAngina_Y',
        'ST_Slope_Down',
        'ST_Slope_Flat',
        'ST_Slope_Up'],
        20)
```

3.4 Preprocessing data: Convert data to numpy arrays

```
[14]: X = np.array(mod_data[features])
      y = np.array(mod_data[output])

      len(features), X.shape, y.shape
```

```
[14]: (20, (918, 20), (918, 1))
```

```
[15]: type(X), type(y), X.dtype, y.dtype
```

```
[15]: (numpy.ndarray, numpy.ndarray, dtype('float64'), dtype('int64'))
```

3.5 Preprocessing data: Splitting data into train and test samples/labels

```
[15]: SEED = 2

      # split dataframe into train samples/labels and test samples/labels
      train_samples, test_samples, train_labels, test_labels = train_test_split(X, y,
      ↪ test_size=0.20, random_state=SEED)
      # convert the dataframes to numpy arrays (tensors)
      train_labels = np.array(train_labels)
      train_samples = np.array(train_samples)
      test_labels = np.array(test_labels)
```

```
test_samples = np.array(test_samples)

train_samples.shape, train_labels.shape, len(features)
```

[15]: ((734, 20), (734, 1), 20)

4 Preprocessing data: Normalizing (z-scoring) data

```
[16]: # converts the numerical values to all be within the range of [0, 1]
scaler = MinMaxScaler(feature_range=[0,1])
# fit and transform train data
scaled_train_samples = scaler.fit_transform(train_samples)
# transform but not fit test data to prevent test data bias leakage
scaled_test_samples = scaler.transform(test_samples)
# verify the data was transformed
scaled_train_samples, scaled_train_samples.shape, scaled_test_samples, ↵
↵scaled_test_samples.shape
```

```
[16]: (array([[0.44897959, 0.8      , 0.      , ..., 0.      , 1.      ,
0.      ],
[0.12244898, 0.75      , 0.35489221, ..., 0.      , 0.      ,
1.      ],
[0.67346939, 0.705     , 0.48424544, ..., 0.      , 1.      ,
0.      ],
...,
[0.46938776, 0.685     , 0.56218905, ..., 0.      , 1.      ,
0.      ],
[0.67346939, 0.695     , 0.46932007, ..., 0.      , 0.      ,
1.      ],
[0.6122449 , 0.675     , 0.3681592 , ..., 0.      , 0.      ,
1.      ]]),
(734, 20),
array([[0.48979592, 0.64      , 0.      , ..., 0.      , 0.      ,
1.      ],
[0.30612245, 0.575     , 0.      , ..., 0.      , 1.      ,
0.      ],
[0.87755102, 0.65      , 0.36650083, ..., 0.      , 1.      ,
0.      ],
...,
[0.71428571, 0.65      , 0.4212272 , ..., 0.      , 1.      ,
0.      ],
[0.67346939, 0.65      , 0.      , ..., 0.      , 1.      ,
0.      ],
[0.59183673, 0.6      , 0.58706468, ..., 0.      , 0.      ,
1.      ]]),
(184, 20))
```

5 Build the Neural Network Model with Keras Sequential class

```
[17]: # building a Sequential model - linear stack of layers
# init the model
model = Sequential()
# add Dense layers to the models
# units = number of nodes, input_shape = tensor shape the input layer expects
# (units weights); activation - activation function
model.add(Dense(units=16, activation='relu', input_shape=(20, )))
model.add(Dense(units=16, activation='relu'))
# use sigmoid in last layer because it is a binary classification problem
model.add(Dense(units=1, activation='sigmoid'))
```

5.1 Visualization of Neural Network Dense Layers

```
[18]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	336
dense_1 (Dense)	(None, 16)	272
dense_2 (Dense)	(None, 1)	17

Total params: 625

Trainable params: 625

Non-trainable params: 0

6 Compile the Model

- Assign loss function: the function that is minimized by the optimizer
- Assign optimizer function: how the model learns and minimizes the loss function
- Choose metrics: used to evaluate the performance of the model

```
[19]: # compilation step - 1) loss function 2) optimizer 3) Metrics to monitor during
# training and testing
model.compile(optimizer=optimizers.RMSprop(learning_rate=0.001), loss=losses.
# binary_crossentropy, metrics=["accuracy"])

# alt ways to construct optimizer
# opt = Adam(learning_rate=0.01) # defining the optimizer
# model.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])
```

```
# model.compile(optimizer='rmsprop', loss='binary_crossentropy',  
↳metrics=['accuracy'])  
# model.compile(optimizer=optimizers.RMSprop(learning_rate=0.001),  
↳loss='binary_crossentropy', metrics=['accuracy'])
```

7 Train the model

- Assign the train samples as x with their train labels as y
- Assign batch size, the number of samples that will be propagated through the network
- Assign epochs, the number of iterations the model will run through the layers

```
[20]: # training the model by fitting the normalized training data  
history = model.fit(x=scaled_train_samples, y=train_labels, batch_size=20,  
↳epochs=30)  
history
```

```
Epoch 1/30  
37/37 [=====] - 0s 703us/step - loss: 0.6031 -  
accuracy: 0.7520  
Epoch 2/30  
37/37 [=====] - 0s 836us/step - loss: 0.4909 -  
accuracy: 0.8270  
Epoch 3/30  
37/37 [=====] - 0s 738us/step - loss: 0.4123 -  
accuracy: 0.8420  
Epoch 4/30  
37/37 [=====] - 0s 916us/step - loss: 0.3705 -  
accuracy: 0.8515  
Epoch 5/30  
37/37 [=====] - 0s 819us/step - loss: 0.3537 -  
accuracy: 0.8529  
Epoch 6/30  
37/37 [=====] - 0s 840us/step - loss: 0.3438 -  
accuracy: 0.8569  
Epoch 7/30  
37/37 [=====] - 0s 976us/step - loss: 0.3373 -  
accuracy: 0.8638  
Epoch 8/30  
37/37 [=====] - 0s 1ms/step - loss: 0.3324 - accuracy:  
0.8624  
Epoch 9/30  
37/37 [=====] - 0s 745us/step - loss: 0.3292 -  
accuracy: 0.8610  
Epoch 10/30  
37/37 [=====] - 0s 804us/step - loss: 0.3283 -  
accuracy: 0.8624  
Epoch 11/30
```


37/37 [=====] - 0s 960us/step - loss: 0.3260 -
 accuracy: 0.8665
 Epoch 12/30
 37/37 [=====] - 0s 833us/step - loss: 0.3231 -
 accuracy: 0.8610
 Epoch 13/30
 37/37 [=====] - 0s 806us/step - loss: 0.3228 -
 accuracy: 0.8692
 Epoch 14/30
 37/37 [=====] - 0s 915us/step - loss: 0.3184 -
 accuracy: 0.8665
 Epoch 15/30
 37/37 [=====] - 0s 883us/step - loss: 0.3206 -
 accuracy: 0.8678
 Epoch 16/30
 37/37 [=====] - 0s 837us/step - loss: 0.3206 -
 accuracy: 0.8665
 Epoch 17/30
 37/37 [=====] - 0s 955us/step - loss: 0.3174 -
 accuracy: 0.8651
 Epoch 18/30
 37/37 [=====] - 0s 916us/step - loss: 0.3173 -
 accuracy: 0.8651
 Epoch 19/30
 37/37 [=====] - 0s 893us/step - loss: 0.3165 -
 accuracy: 0.8665
 Epoch 20/30
 37/37 [=====] - 0s 881us/step - loss: 0.3164 -
 accuracy: 0.8665
 Epoch 21/30
 37/37 [=====] - 0s 1ms/step - loss: 0.3143 - accuracy:
 0.8665
 Epoch 22/30
 37/37 [=====] - 0s 1ms/step - loss: 0.3133 - accuracy:
 0.8651
 Epoch 23/30
 37/37 [=====] - 0s 928us/step - loss: 0.3113 -
 accuracy: 0.8665
 Epoch 24/30
 37/37 [=====] - 0s 917us/step - loss: 0.3095 -
 accuracy: 0.8719
 Epoch 25/30
 37/37 [=====] - 0s 1ms/step - loss: 0.3114 - accuracy:
 0.8760
 Epoch 26/30
 37/37 [=====] - 0s 968us/step - loss: 0.3079 -
 accuracy: 0.8665
 Epoch 27/30

```

37/37 [=====] - 0s 761us/step - loss: 0.3082 -
accuracy: 0.8706
Epoch 28/30
37/37 [=====] - 0s 874us/step - loss: 0.3081 -
accuracy: 0.8665
Epoch 29/30
37/37 [=====] - 0s 1ms/step - loss: 0.3068 - accuracy:
0.8706
Epoch 30/30
37/37 [=====] - 0s 931us/step - loss: 0.3060 -
accuracy: 0.8733

```

[20]: <keras.callbacks.History at 0x7fa6baecc100>

8 Evaluate the Model's performance

- The evaluate method takes in a test sample numpy array and their associated test labels
- Returns the loss value & metrics value (accuracy score) for the model in test mode
- Loss is the scalar value that is attempted to be minimized during training of the model. The lower the loss, the closer our predictions are to the true labels.

```

[21]: results = model.evaluate(scaled_test_samples, test_labels)
      results

```

```

6/6 [=====] - 0s 742us/step - loss: 0.3765 - accuracy:
0.8315

```

[21]: [0.37645289301872253, 0.83152174949646]

9 Further Experiments

-

9.1 Hidden Layers

- Try using one or three hidden layers, and see how doing so affects validation and test accuracy.

-

9.2 Hidden Units

- Try using layers with more hidden units or fewer hidden units: 32 units, 64 units, and so on.

-

9.3 Loss Functions

- Try using the mse loss function instead of binary_crossentropy.

•

9.4 Activation Function

- Try using the tanh activation (an activation that was popular in the early days of neural networks)

9.5 Reusable function for creating Keras Sequential Models

```
[17]: DEFAULT_LAYERS = [  
        Dense(units=16, activation='relu', input_shape=(20, )),  
        Dense(units=32, activation='relu'),  
        Dense(units=1, activation='sigmoid')  
    ]  
DEFAULT_OPTIMIZER = optimizers.RMSprop(learning_rate=0.001)  
DEFAULT_LOSS = losses.binary_crossentropy  
DEFAULT_METRICS = ["accuracy"]  
  
[18]: def create_neural_network(scaled_train_samples,  
                                train_labels,  
                                layers = DEFAULT_LAYERS,  
                                optimizer=DEFAULT_OPTIMIZER,  
                                loss=DEFAULT_LOSS,  
                                metrics=DEFAULT_METRICS):  
    model = Sequential(layers)  
    model.compile(optimizer=optimizer, loss=loss, metrics=metrics)  
    model.fit(x=scaled_train_samples, y=train_labels, batch_size=20, epochs=30,  
↳ verbose=0)  
    return model  
  
def get_nn_results(model, scaled_test_samples, test_labels):  
    return model.evaluate(scaled_test_samples, test_labels)
```

9.6 Further Experiments: Hidden Layers

```
[24]: layers_dict = {  
    # 1 hidden layer  
    '1': [  
        Dense(units=1, activation='sigmoid')  
    ],  
    # 2 hidden layers  
    '2': [  
        Dense(units=16, activation='relu', input_shape=(20, )),  
        Dense(units=1, activation='sigmoid')  
    ]  
}
```

```

    ],
    # 3 hidden layers
    '3': [
        Dense(units=16, activation='relu', input_shape=(20, )),
        Dense(units=16, activation='relu'),
        Dense(units=1, activation='sigmoid')
    ]
}
results_layers_dict = {}

```

```

[25]: SEED = 2

train_samples, test_samples, train_labels, test_labels = train_test_split(X, y,
    ↳test_size=0.20, random_state=SEED)
# convert the dataframes to numpy arrays (tensors)
train_labels = np.array(train_labels)
train_samples = np.array(train_samples)
test_labels = np.array(test_labels)
test_samples = np.array(test_samples)

# converts the numerical values to all be within the range of [0, 1]
scaler = MinMaxScaler(feature_range=[0,1])
# fit and transform train data
scaled_train_samples = scaler.fit_transform(train_samples)
# transform but not fit test data to prevent test data bias leakage
scaled_test_samples = scaler.transform(test_samples)
# verify the data was transformed
scaled_train_samples, scaled_train_samples.shape, scaled_test_samples,
    ↳scaled_test_samples.shape

for key, layers in layers_dict.items():
    curr_model = create_neural_network(scaled_train_samples, train_labels,
    ↳layers)
    results_layers_dict[key] = get_nn_results(curr_model, scaled_test_samples,
    ↳test_labels)

print('\n')
for key, value in results_layers_dict.items():
    print(f'Achieved {value[0]:.4f} loss {value[1]:.4f} accuracy with {key},
    ↳layer(s)')

```

```

6/6 [=====] - 0s 1ms/step - loss: 0.3875 - accuracy:
0.8641
6/6 [=====] - 0s 940us/step - loss: 0.3684 - accuracy:
0.8696
6/6 [=====] - 0s 968us/step - loss: 0.3511 - accuracy:
0.8587

```

Achieved 0.3875 loss 0.8641 accuracy with 1 layer(s)
Achieved 0.3684 loss 0.8696 accuracy with 2 layer(s)
Achieved 0.3511 loss 0.8587 accuracy with 3 layer(s)

9.7 Further Experiments: Hidden Units

```
[26]: units_dict = {  
    # 16 hidden layer  
    '16': [  
        Dense(units=16, activation='relu', input_shape=(20, )),  
        Dense(units=16, activation='relu'),  
        Dense(units=1, activation='sigmoid')  
    ],  
    # 32 hidden layers  
    '32': [  
        Dense(units=32, activation='relu', input_shape=(20, )),  
        Dense(units=32, activation='relu'),  
        Dense(units=1, activation='sigmoid')  
    ],  
    # 64 hidden layers  
    '64': [  
        Dense(units=64, activation='relu', input_shape=(20, )),  
        Dense(units=64, activation='relu'),  
        Dense(units=1, activation='sigmoid')  
    ]  
}  
results_units_dict = {}
```

```
[27]: SEED = 2  
  
train_samples, test_samples, train_labels, test_labels = train_test_split(X, y,  
    ↪ test_size=0.20, random_state=SEED)  
# convert the dataframes to numpy arrays (tensors)  
train_labels = np.array(train_labels)  
train_samples = np.array(train_samples)  
test_labels = np.array(test_labels)  
test_samples = np.array(test_samples)  
  
# converts the numerical values to all be within the range of [0, 1]  
scaler = MinMaxScaler(feature_range=[0,1])  
# fit and transform train data  
scaled_train_samples = scaler.fit_transform(train_samples)  
# transform but not fit test data to prevent test data bias leakage  
scaled_test_samples = scaler.transform(test_samples)  
# verify the data was transformed
```

```

scaled_train_samples, scaled_train_samples.shape, scaled_test_samples,
↳scaled_test_samples.shape

for key, layers in units_dict.items():
    curr_model = create_neural_network(scaled_train_samples, train_labels,
↳layers)
    results_units_dict[key] = get_nn_results(curr_model, scaled_test_samples,
↳test_labels)

print('\n')
for key, value in results_units_dict.items():
    print(f'Achieved {value[0]:.4f} loss {value[1]:.4f} accuracy with {key}
↳units/nodes per deep layer')

```

```

6/6 [=====] - 0s 1ms/step - loss: 0.3605 - accuracy:
0.8587
6/6 [=====] - 0s 816us/step - loss: 0.3748 - accuracy:
0.8370
6/6 [=====] - 0s 846us/step - loss: 0.3888 - accuracy:
0.8370

```

Achieved 0.3605 loss 0.8587 accuracy with 16 units/nodes per deep layer
 Achieved 0.3748 loss 0.8370 accuracy with 32 units/nodes per deep layer
 Achieved 0.3888 loss 0.8370 accuracy with 64 units/nodes per deep layer

9.8 Further Experiments: Loss Functions

```

[28]: loss_functions = ['binary_crossentropy', 'hinge', 'squared_hinge']
      results_loss_dict = {}

```

```

[29]: SEED = 2

train_samples, test_samples, train_labels, test_labels = train_test_split(X, y,
↳test_size=0.20, random_state=SEED)
# convert the dataframes to numpy arrays (tensors)
train_labels = np.array(train_labels)
train_samples = np.array(train_samples)
test_labels = np.array(test_labels)
test_samples = np.array(test_samples)

# converts the numerical values to all be within the range of [0, 1]
scaler = MinMaxScaler(feature_range=[0,1])
# fit and transform train data
scaled_train_samples = scaler.fit_transform(train_samples)
# transform but not fit test data to prevent test data bias leakage

```

```

scaled_test_samples = scaler.transform(test_samples)
# verify the data was transformed
scaled_train_samples, scaled_train_samples.shape, scaled_test_samples,
    ↪scaled_test_samples.shape

for loss_func in loss_functions:
    curr_model = create_neural_network(scaled_train_samples, train_labels,
    ↪loss=loss_func)
    results_loss_dict[loss_func] = get_nn_results(curr_model,
    ↪scaled_test_samples, test_labels)

print('\n')
for key, value in results_loss_dict.items():
    print(f'Achieved {value[0]:.4f} loss {value[1]:.4f} accuracy with the {key}
    ↪loss function')

```

```

6/6 [=====] - 0s 854us/step - loss: 0.3638 - accuracy:
0.8696
6/6 [=====] - 0s 916us/step - loss: 0.6380 - accuracy:
0.8478
6/6 [=====] - 0s 739us/step - loss: 0.7288 - accuracy:
0.8370

```

Achieved 0.3638 loss 0.8696 accuracy with the binary_crossentropy loss function
 Achieved 0.6380 loss 0.8478 accuracy with the hinge loss function
 Achieved 0.7288 loss 0.8370 accuracy with the squared_hinge loss function

9.9 Changing feature range to be [-1,1] for tanh activation

```

[21]: SEED = 2
tanh_loss_dict = {}
train_samples, test_samples, train_labels, test_labels = train_test_split(X, y,
    ↪test_size=0.20, random_state=SEED)
# convert the dataframes to numpy arrays (tensors)
train_labels = np.array(train_labels)
train_samples = np.array(train_samples)
test_labels = np.array(test_labels)
test_samples = np.array(test_samples)
layers = [
    Dense(units=16, activation='relu', input_shape=(20, )),
    Dense(units=32, activation='relu'),
    Dense(units=1, activation='tanh')
]

# converts the numerical values to all be within the range of [-1, 1]
scaler = MinMaxScaler(feature_range=[-1,1])

```

```

# fit and transform train data
scaled_train_samples = scaler.fit_transform(train_samples)
# transform but not fit test data to prevent test data bias leakage
scaled_test_samples = scaler.transform(test_samples)
# verify the data was transformed
scaled_train_samples, scaled_train_samples.shape, scaled_test_samples,
↳scaled_test_samples.shape

curr_model = create_neural_network(scaled_train_samples, train_labels,
↳loss=losses.binary_crossentropy)
tanh_loss_dict['tanh'] = get_nn_results(curr_model, scaled_test_samples,
↳test_labels)

tanh_loss_dict

```

6/6 [=====] - 0s 843us/step - loss: 0.4098 - accuracy: 0.8533

[21]: {'tanh': [0.40980401635169983, 0.85326087474823]}

9.10 Further Experiments: Activation Functions

```

[30]: activation_functions_dict = {
    # sigmoid
    'sigmoid': [
        Dense(units=16, activation='relu', input_shape=(20, )),
        Dense(units=16, activation='relu'),
        Dense(units=1, activation='sigmoid')
    ],
    # tanh
    'tanh': [
        Dense(units=16, activation='relu', input_shape=(20, )),
        Dense(units=16, activation='relu'),
        Dense(units=1, activation='tanh')
    ],
    # relu
    'relu': [
        Dense(units=16, activation='relu', input_shape=(20, )),
        Dense(units=16, activation='relu'),
        Dense(units=1, activation='relu')
    ]
}
results_activation_functions_dict = {}

```

[31]: SEED = 2


```

train_samples, test_samples, train_labels, test_labels = train_test_split(X, y,
    ↳test_size=0.20, random_state=SEED)
# convert the dataframes to numpy arrays (tensors)
train_labels = np.array(train_labels)
train_samples = np.array(train_samples)
test_labels = np.array(test_labels)
test_samples = np.array(test_samples)

# converts the numerical values to all be within the range of [0, 1]
scaler = MinMaxScaler(feature_range=[0,1])
# fit and transform train data
scaled_train_samples = scaler.fit_transform(train_samples)
# transform but not fit test data to prevent test data bias leakage
scaled_test_samples = scaler.transform(test_samples)
# verify the data was transformed
scaled_train_samples, scaled_train_samples.shape, scaled_test_samples,
    ↳scaled_test_samples.shape

for key, layers in activation_functions_dict.items():
    curr_model = create_neural_network(scaled_train_samples, train_labels,
    ↳layers)
    results_activation_functions_dict[key] = get_nn_results(curr_model,
    ↳scaled_test_samples, test_labels)

print('\n')
for key, value in results_activation_functions_dict.items():
    print(f'Achieved {value[0]:.4f} loss {value[1]:.4f} accuracy with the {key}
    ↳activation function.')

```

```

6/6 [=====] - 0s 792us/step - loss: 0.3668 - accuracy:
0.8533
6/6 [=====] - 0s 721us/step - loss: 0.4986 - accuracy:
0.8207
6/6 [=====] - 0s 1ms/step - loss: 0.7425 - accuracy:
0.8370

```

Achieved 0.3668 loss 0.8533 accuracy with the sigmoid activation function.
 Achieved 0.4986 loss 0.8207 accuracy with the tanh activation function.
 Achieved 0.7425 loss 0.8370 accuracy with the relu activation function.

[]: